AI Recruiter

Role	Senior AI Engineer
Company	Redrob AI
Date	24 June 2026
Shortlist	30 candidates
Pool sized	30 candidates scored

Generated by AI Recruiter Pipeline — Evidence-based ranking using career history, skill depth, behavioral signals, and JD fit analysis.

Executive Summary

Top 5 shortlisted candidates ranked by composite score. Scores are evidence-based and verified against actual profile data.

Rank	Name	Composite	Fit	Impact	Potential	Risk	Key Strength
1	Ira Dalal	92.25	92	95	94	15	Strong production experience with hybrid retriev
2	Mira Ghosh	91.05	92	95	88	15	Deep production experience with hybrid search a
3	Ishaan Arora	91.05	92	95	88	15	Proven track record of shipping production ML s
4	Vivaan Shah	89.6	90	95	88	20	Strong product-company background
5	Aarav Agarwal	89.55	92	95	88	25	Strong production-first mindset ('shipping in 6 we

One-line assessment:

1. Ira Dalal — Ira Dalal earned the top ranking due to her exceptional blend of technical depth and measurable business impact. Her hands-on experience designing hybrid retrieval systems (BM25 + dense vector) demonstrates…

2. Mira Ghosh — Mira Ghosh earns the #2 ranking due to her exceptional technical execution and measurable infrastructure impact. Her ability to fine-tune LLaMA-2-7B and Mistral-7B using LoRA while simultaneously engineering a 200K…

3. Ishaan Arora — Ishaan Arora earned a #3 ranking due to his exceptional technical execution and proven ability to drive product metrics at scale. His work building a recommendation system for 10M+ users,…

4. Vivaan Shah — Vivaan Shah earns a #4 ranking due to his exceptional technical proficiency and proven ability to drive bottom-line results. His deep expertise in LTR and embedding-based retrieval, combined with a…

5. Aarav Agarwal — Aarav Agarwal earns a top-five ranking due to his exceptional technical depth and proven ability to deliver high-impact results. His hands-on experience designing a hybrid retrieval system (BM25 + dense…

#1 Ira Dalal [DARK HORSE]

Composite	Fit	Impact	Potential	Risk	Confidence
92.25	92	95	94	15	high

GREEN FLAGS

• Strong production experience with hybrid retrieval and vector database
• Quantified business impact (retention, cost, latency)
• Pragmatic 'ship-first' product engineering mindset
• Stable career history at reputable companies

YELLOW FLAGS

• Currently based in Trivandrum; relocation to Pune/Noida required
• Experience is heavily weighted toward large-scale consumer tech; needs

SKILL GAPS

None identified

Interview Questions

1. You successfully scaled a hybrid retrieval system for 30M+ profiles; walk me through the specific trade-offs you faced when balancing the recall performance of your dense vector em

2. While your technical implementation is strong, I’d like to understand your approach to the 'Evaluation' side of ranking systems: what specific offline metrics and online guardrails

3. You’ve emphasized a preference for shipping working systems in 6 weeks over theoretical perfection; describe a time when this 'ship-first' approach led to a significant technical d

Ranking Rationale

Ira Dalal earned the top ranking due to her exceptional blend of technical depth and measurable business impact. Her hands-on experience designing hybrid retrieval systems (BM25 + dense vector) demonstrates high-level engineering proficiency, while her ability to move a recommendation system from offline to A/B testing in just five weeks highlights a pragmatic, high-velocity mindset. Most impressively, Ira Dalal delivered tangible results, specifically driving an 11% increase in new-user retention through Thompson sampling. Her combination of specialized vector database expertise…

#2 Mira Ghosh [DARK HORSE]

Composite	Fit	Impact	Potential	Risk	Confidence
91.05	92	95	88	15	high

GREEN FLAGS

• Deep production experience with hybrid search and ranking evaluation
• Quantifiable impact on latency and infrastructure costs
• Strong alignment with 'systems over frameworks' philosophy
• Experience in both startup and large-scale product environments

YELLOW FLAGS

• Location is Chennai, while JD prefers Pune/Noida/Hyderabad/Mumbai/Delh
• GitHub activity score is -1 (no public open-source footprint visible)

SKILL GAPS

None identified

Interview Questions

1. You mentioned building a data curation pipeline that generated 200K preference pairs; walk me through the specific quality control mechanisms or heuristics you implemented to ensur

2. While your experience with vector databases and retrieval is strong, could you describe a scenario where you had to debug a significant performance degradation in a hybrid search s

3. You hold a strong view that ranking problems are solved by feature engineering and evaluation rather than model size; tell me about a time you pushed back against a stakeholder or

Ranking Rationale

Mira Ghosh earns the #2 ranking due to her exceptional technical execution and measurable infrastructure impact. Her ability to fine-tune LLaMA-2-7B and Mistral-7B using LoRA while simultaneously engineering a 200K preference pair data pipeline demonstrates rare end-to-end proficiency. Furthermore, her deployment of these models via BentoML on Kubernetes, achieving a sub-200ms p95 latency, directly validates her 'systems over frameworks' philosophy. Mira Ghosh combines high-level model optimization with rigorous production-grade engineering, making her a high-impact asset capable of delivering scalable, performant…

#3 Ishaan Arora [DARK HORSE]

Composite	Fit	Impact	Potential	Risk	Confidence
91.05	92	95	88	15	high

GREEN FLAGS

• Proven track record of shipping production ML systems at scale (10M+ u
• Strong alignment with 'product-first' engineering culture
• Deep expertise in the specific stack requested (Embeddings, Vector DBs
• Long-term tenure at previous companies (36 and 38 months)

YELLOW FLAGS

• Current location (Ahmedabad) is not in the preferred list of cities
• Notice period of 45 days is standard but requires coordination

SKILL GAPS

None identified

Interview Questions

1. You achieved a 35% improvement in search relevance over BM25 by migrating to embedding-based retrieval; walk me through the specific trade-offs you encountered between latency and

2. While you have extensive experience with vector databases and retrieval, we need to ensure deep rigor in evaluation; can you describe a time you had to design a custom evaluation f

3. You mentioned that your primary learning comes from shipping systems and seeing what breaks under production load; tell me about a time when a model that performed well in offline

Ranking Rationale

Ishaan Arora earned a #3 ranking due to his exceptional technical execution and proven ability to drive product metrics at scale. His work building a recommendation system for 10M+ users, which directly improved 7-day retention by 6%, demonstrates high impact and technical maturity. Ishaan Arora possesses deep, specialized expertise in the required stack—specifically embeddings, vector databases, and ranking—and aligns perfectly with a product-first engineering culture. While his technical output is elite, he is positioned at #3 behind individuals with slightly…

#4 Vivaan Shah

Composite	Fit	Impact	Potential	Risk	Confidence
89.6	90	95	88	20	high

GREEN FLAGS
• Strong product-company background
• Quantified production impact
• Deep experience with LTR and embedding-based retrieval
• Strong focus on evaluation and offline-to-online correlation

YELLOW FLAGS
• Location is Kolkata (requires relocation to Pune/Noida)
• 90-day notice period

SKILL GAPS
None identified

Interview Questions

1. You mentioned evolving the ranking layer at your previous role from a hand-tuned scoring function to a learning-to-rank model; walk me through the specific architecture you chose,

2. While you have extensive experience with vector databases like Qdrant and FAISS, our infrastructure requires complex hybrid search implementations; can you describe a scenario wher

3. Given your focus on shipping systems that hold up under production load, tell me about a time when a model you deployed performed well in offline benchmarks but failed to deliver t

Ranking Rationale

Vivaan Shah earns a #4 ranking due to his exceptional technical proficiency and proven ability to drive bottom-line results. His deep expertise in LTR and embedding-based retrieval, combined with a robust stack including Qdrant, FAISS, and MLOps, positions him as a high-impact contributor. Most notably, Vivaan Shah demonstrated significant business value at Freshworks by improving revenue-per-search by 12%, alongside measurable gains in 7-day retention and session duration. His strong product-company background and lack of identified skill gaps make him a…

#5 Aarav Agarwal [DARK HORSE]

Composite	Fit	Impact	Potential	Risk	Confidence
89.55	92	95	88	25	high

GREEN FLAGS
• Strong production-first mindset ('shipping in 6 weeks')
• Quantified impact across multiple roles
• Deep technical expertise in hybrid search and ranking
• Experience mentoring junior engineers

YELLOW FLAGS
• Location mismatch (Kochi vs. preferred hubs)
• One short tenure (6 months at Aganitha)
• Notice period of 60 days

SKILL GAPS
None identified

Interview Questions

1. You achieved a 60% reduction in p95 latency while simultaneously improving NDCG@10 by 18% on your hybrid retrieval system; walk me through the specific architectural trade-offs you

2. While you have extensive experience in building and deploying retrieval systems, our role requires a heavy emphasis on MLOps—specifically regarding automated model monitoring, drif

3. You mentioned a preference for shipping working systems in 6 weeks over theoretical perfection; tell me about a time when this 'fast-shipping' approach led to a significant technic

Ranking Rationale

Aarav Agarwal earns a top-five ranking due to his exceptional technical depth and proven ability to deliver high-impact results. His hands-on experience designing a hybrid retrieval system (BM25 + dense vector) and implementing exploration-exploitation policies via Thompson sampling demonstrates advanced engineering maturity. Aarav Agarwal’s production-first mindset, evidenced by his ability to ship complex features in six weeks, combined with his specialized proficiency in QLoRA, pgvector, and GANs, makes him a high-impact asset. With no identified skill gaps and a low…

#6 Naina Tiwari [DARK HORSE]

Composite	Fit	Impact	Potential	Risk	Confidence
87.1	85	95	88	25	high

GREEN FLAGS

• Proven track record of shipping production ML systems at scale
• Strong focus on data quality and infrastructure over just modeling
• Quantifiable business impact across all previous roles
• Experience with both search and recommendation systems

YELLOW FLAGS

• Currently based in the USA (requires relocation to India)
• No explicit evidence of mentoring or team leadership
• Relatively short tenure at Zoho (13 months)

SKILL GAPS

• Mentorship/Team leadership experience (not explicitly evidenced)

Interview Questions

1. You achieved a 12% increase in revenue-per-search at Netflix; walk me through the specific architectural trade-offs you made between latency and model complexity to achieve that re

2. While your technical background in search infrastructure is strong, our team requires someone who can mentor junior engineers and guide technical strategy; can you describe a time

3. You've noted that you prioritize shipping real systems over academic theory and focus heavily on data quality; tell me about a time when a production search system you built failed

Ranking Rationale

Naina Tiwari ranks 6th due to her exceptional technical execution and proven ability to drive high-impact business outcomes. Her track record is highlighted by a 12% increase in revenue-per-search at Netflix and a 35% improvement in search-relevance at Niramai, demonstrating her proficiency with Kubeflow, FAISS, and complex feature engineering. While her technical infrastructure expertise is elite, she ranks 6th because her application lacks explicit evidence of mentorship or team leadership experience, which is a critical requirement for higher-level roles. Her…

#7 Sunil Mishra

Composite	Fit	Impact	Potential	Risk	Confidence
86.6	90	85	88	20	high

GREEN FLAGS

• Quantified business impact (12% revenue lift, 35% relevance improvemen
• Deep production experience with vector search and ranking
• Strong product-engineering mindset (collaborating with PMs on optimiza
• Mentorship experience

YELLOW FLAGS

• Currently based in Kolkata; requires relocation to Pune/Noida
• GitHub activity score is -1 (no public code visibility)

SKILL GAPS

None identified

Interview Questions

1. You achieved a 35% improvement in search relevance over BM25; walk me through the specific evaluation framework you used to validate this, and how you handled the trade-offs betwee

2. While your experience with traditional ML models like XGBoost is clear, how have you specifically adapted your retrieval and ranking pipelines to incorporate Large Language Models

3. You mentioned that defining the optimization target with PMs was as important as the modeling itself; tell me about a time when the initial metric you chose didn't align with the d

Ranking Rationale

Sunil Mishra earns a strong ranking due to his proven ability to bridge the gap between complex ML infrastructure and tangible business outcomes. His technical proficiency in building semantic search systems using sentence-transformers is validated by a measurable 12% revenue lift and 35% improvement in search relevance. Furthermore, Sunil Mishra demonstrates a sophisticated product-engineering mindset, evidenced by his successful collaboration with PMs to define optimization targets and his operational expertise in managing production pipelines with MLflow and Kubernetes. He is…

#8 Suresh Dutta [DARK HORSE]

Composite	Fit	Impact	Potential	Risk	Confidence
86.6	90	85	88	20	high

GREEN FLAGS
• Quantified impact on production systems
• Strong balance of foundational ML and modern LLM/RAG experience
• Experience with evaluation frameworks and MLOps
• Mentorship experience

YELLOW FLAGS
• Currently based in Chennai (requires relocation to Pune/Noida/etc.)
• 90-day notice period is on the higher side for a startup

SKILL GAPS
None identified

Interview Questions

1. You achieved a 35% improvement in search relevance over BM25; walk me through the specific evaluation framework you used to validate these human relevance judgments, and how you ha

2. While you have extensive experience with vector databases and RAG, our role requires deep expertise in hybrid search infrastructure; can you describe a scenario where you had to ba

3. You mentioned that your learning comes from seeing systems hold up under production load; tell me about a time when a production ML pipeline failed or underperformed in a way you d

Ranking Rationale

Suresh Dutta earns a high ranking due to his proven ability to deliver measurable technical impact. His development of a semantic search feature for 500,000 documents, which improved relevance by 35% over Elasticsearch BM25, demonstrates exceptional engineering proficiency. Furthermore, Suresh Dutta shows a strong command of the modern AI stack, evidenced by his successful implementation of RAG-based chatbots using OpenAI embeddings and Pinecone. His balanced expertise in foundational ML, MLOps, and evaluation frameworks makes him a high-impact contributor with significant…

#9 Meera Chowdary [DARK HORSE]

Composite	Fit	Impact	Potential	Risk	Confidence
86.6	90	85	88	20	high

GREEN FLAGS

• Quantified impact on business metrics (retention, session time, releva
• Deep production experience with vector search and ranking
• Strong product-company background (Dream11, Krutrim)
• Pragmatic 'evaluation-first' mindset

YELLOW FLAGS

• Current location is Chennai; role requires Pune/Noida/Hyderabad/Mumbai
• Not currently 'open to work' (passive candidate)

SKILL GAPS

None identified

Interview Questions

1. You achieved a 35% improvement in search relevance over BM25; walk me through the specific evaluation methodology you used to validate this, and how you isolated the impact of the

2. While you have extensive experience in building and deploying models, our role requires a heavy focus on MLOps; can you describe a time you had to move a model from a prototype to

3. You mentioned that most retrieval problems are actually evaluation problems in disguise; tell me about a time when you realized your initial evaluation metrics were misaligned with

Ranking Rationale

Meera Chowdary earned a top-tier ranking due to her exceptional technical depth in vector search and production-grade recommendation systems. Her proven ability to scale semantic search for 500,000 documents and manage content recommendations for over 10 million users at companies like Dream11 and Krutrim demonstrates high-impact engineering capability. With deep expertise in OpenSearch, Elasticsearch, and Haystack, Meera Chowdary shows no significant skill gaps. Her combination of quantified business results—specifically in retention and session time—and robust experience with complex data architectures…

#10 Anika Rao

Composite	Fit	Impact	Potential	Risk	Confidence
86.35	85	95	88	30	high

GREEN FLAGS

• Proven track record of shipping production ML systems at scale
• Deep expertise in both traditional IR (BM25) and modern vector/hybrid
• Strong focus on evaluation frameworks (NDCG, MRR, A/B testing)
• Quantifiable impact on latency and model performance

YELLOW FLAGS

• Currently based in London (requires relocation to India)
• GitHub activity score is low (-1), which contradicts the 'open-source

SKILL GAPS

• None identified; candidate meets or exceeds all must-have technical re

Interview Questions

1. You achieved a 60% reduction in p95 latency while simultaneously improving NDCG@10 by 18% during your migration to a learning-to-rank model; walk me through the specific architectu

2. While your background is strong in model development and offline evaluation, our role requires significant ownership of MLOps pipelines; can you describe a time you had to troubles

3. You prioritize shipping a working system in 6 weeks over a perfect one in 6 months; tell me about a time when 'shipping early' forced you to make a technical compromise that you la

Ranking Rationale

Anika Rao is ranked #10 due to her exceptional technical proficiency and high-impact delivery. Her leadership in migrating to learning-to-rank systems and deploying semantic search for over 35 million items demonstrates a rare ability to scale complex ML infrastructure. Anika Rao’s pragmatic approach—prioritizing functional, production-ready systems within six-week cycles—aligns perfectly with our operational goals. With deep expertise in both traditional IR and modern vector search, combined with a rigorous focus on evaluation metrics like NDCG and MRR, she represents a…

Dark Horse Candidates

These candidates ranked below position 15 in vector similarity search but show high impact or potential scores (>= 75) with fit >= 50. A traditional ATS would miss them.

Ira Dalal (Rank 1 · Composite 92.25)

Ira Dalal is a classic 'hidden gem' who likely ranked lower due to the specific keyword-heavy nature of traditional ATS filters, despite possessing elite-level experience in the exact infrastructure required. Her ability to bridge the gap between complex LLM fine-tuning and production-grade, low-latency retrieval systems is rare. She demonstrates a 'ship-first' engineering maturity that is often missing in candidates who…

Transferable skill mappings:

• Hybrid retrieval system (BM25 + dense vector) for 30M+ profiles → Embeddings-based retrieval systems
• INT8 quantization and batching for sub-200ms p95 latency → MLOps
• Thompson sampling for retention optimization → Evaluation frameworks for ranking systems

---

Mira Ghosh (Rank 2 · Composite 91.05)

Mira is a classic 'hidden gem' who likely ranked lower due to the specific phrasing of her experience (e.g., 'data curation' and 'inference cost optimization') rather than the exact 'Vector Database' keyword matching. Her deep, hands-on experience with the entire lifecycle of retrieval systems—from BM25/dense retrieval to offline evaluation and production deployment—makes her a high-impact candidate who possesses the exact…

Transferable skill mappings:

• BGE, FAISS HNSW, and RAG-based ranking pipelines → Embeddings-based retrieval systems
• Deployed model via BentoML on Kubernetes → MLOps
• Offline evaluation framework (NDCG, MRR, recall@K) → Evaluation frameworks for ranking systems

---

Ishaan Arora (Rank 3 · Composite 91.05)

Ishaan is a classic 'hidden gem' who likely ranked lower due to the specific phrasing of his experience (e.g., 'content recommendation' vs. 'e-commerce search'). His technical stack (Milvus, Sentence Transformers) is a perfect match for the JD, and his ability to drive significant business metrics (6% retention lift) demonstrates a product-first mindset that often outweighs pure academic or keyword-matched candidates.

Transferable skill mappings:

• Content recommendation system for 10M+ users → Embeddings-based retrieval systems
• Semantic search for 500K documents → Vector databases or hybrid search infrastructure
• A/B testing infrastructure and feature engineering → Evaluation frameworks for ranking systems

---

Aarav Agarwal (Rank 5 · Composite 89.55)

Aarav is a classic 'hidden gem' because his profile is heavily weighted toward high-impact engineering outcomes (latency reduction, NDCG improvement) rather than just keyword-matching the JD. While his rank is lower due to the retrieval system's focus on specific terminology, his hands-on experience building production-grade hybrid search systems (BM25 + dense vector) at scale (50M+ queries) makes him a superior…

Transferable skill mappings:

• Designed hybrid retrieval system (BM25 + dense vector) → Embeddings-based retrieval systems / Vector
• Improved NDCG@10 by 18% → Evaluation frameworks for ranking systems
• Reduced p95 retrieval latency by 60% → MLOps

Naina Tiwari (Rank 6 · Composite 87.1)

Naina is a classic 'hidden gem' who likely ranked lower due to the ATS/vector search prioritizing specific job titles or seniority keywords over the sheer scale and diversity of her technical impact. Her experience spans three distinct domains (streaming, health-tech, and SaaS), demonstrating an ability to apply core ML infrastructure principles across vastly different data distributions. She possesses the exact…

Transferable skill mappings:

• Evolved e-commerce search from hand-tuned to learning-to-rank → Evaluation frameworks for ranking sy
• Built content recommendation system serving 10M+ users → Embeddings-based retrieval systems
• Used Kubeflow, BigQuery, and Feature Engineering → MLOps

Suresh Dutta (Rank 8 · Composite 86.6)

Suresh is a classic 'hidden gem' who likely ranked lower due to a lack of specific keyword density in his resume (e.g., perhaps missing specific industry-standard buzzwords for 'evaluation frameworks' despite having the practical experience). His high impact scores (85/88) and low risk score (20) indicate a candidate who has already solved the exact problems the JD requires at scale,…

Transferable skill mappings:

• Developed semantic search feature for 500K documents using sentence-transformers → Embeddings-based
• Improved search-relevance by 35% over Elasticsearch BM25 → Evaluation frameworks for ranking systems
• Built production ML pipelines using MLflow, Kubeflow, and internal feature store → MLOps

---

Meera Chowdary (Rank 9 · Composite 86.6)

Meera is a classic 'hidden gem' who likely ranked lower due to the specific terminology in her resume (e.g., 'collaborative filtering' or 'content recommendation') not perfectly aligning with the 'embeddings-based retrieval' keyword search, despite her having deep, hands-on production experience in exactly those domains. Her 'evaluation-first' philosophy and proven ability to drive significant business metrics (35% relevance improvement) make her…

Transferable skill mappings:

• Collaborative filtering for 10M+ users → Embeddings-based retrieval systems
• Semantic search for 500K documents using sentence-transformers and FAISS → Vector databases or hybri
• Evaluation-first mindset / Search relevance optimization → Evaluation frameworks for ranking systems

---

Om Mishra (Rank 12 · Composite 84.4)

Om Mishra is a classic 'false negative' in automated ranking systems. While his technical stack (Spring Boot, OpenCV) might have diluted his ranking against pure-play Python/ML candidates, his actual experience—specifically shipping production-grade embedding systems for 10M+ users—is superior to many candidates who list the keywords but lack the scale-related engineering maturity. He possesses the rare combination of deep infrastructure knowledge…

Transferable skill mappings:

• Spring Boot / Java-based backend architecture → Python / MLOps
• Semantic search for 500K documents using sentence-transformers and FAISS → Vector databases or hybri
• Offline-to-online correlation and evaluation rigor → Evaluation frameworks for ranking systems

---

Pranav Shah (Rank 14 · Composite 81.4)

Pranav is a classic 'high-impact, low-keyword' candidate. While his resume lacks the specific 'distributed systems' buzzwords often used as filters, his proven track record of shipping production-grade RAG and ranking systems—specifically achieving a 31% reduction in resolution time—demonstrates the practical engineering maturity required for the role. He possesses the exact technical stack (Pinecone, Milvus, PyTorch, Kubeflow) required by the JD,…

Transferable skill mappings:

• Production load testing and deployment of RAG/Ranking systems → Distributed systems optimization
• Evolving search ranking from hand-tuned to ML-based → Evaluation frameworks for ranking systems
• Kubeflow and production monitoring → MLOps

Nisha Bansal (Rank 17 · Composite 81.4)

Nisha is a classic 'hidden gem' who likely ranked lower due to the specific keyword-heavy nature of the retrieval system (e.g., her experience with Flask/Redux or specific CV/Speech libraries might have diluted the vector match for a pure RAG/Search role). However, her actual project experience—specifically migrating from BM25 to embedding-based retrieval—is a perfect functional match for the JD. She possesses…

Transferable skill mappings:

• Migrating keyword-search (BM25) to embedding-based retrieval → Embeddings-based retrieval systems
• Operational experience with drift, A/B testing, and retraining → Evaluation frameworks for ranking s
• Built RAG-based chatbot using Pinecone and OpenAI embeddings → Vector databases or hybrid search inf

---

Priya Bhatia (Rank 22 · Composite 76.9)

Priya is a classic 'hidden gem' who likely ranked lower due to the breadth of her technical stack (GANs, TTS, Forecasting) diluting the keyword-matching signal for a pure search/retrieval role. Her experience migrating a legacy keyword-search system to an embedding-based architecture is a high-value, rare 'been-there-done-that' experience that directly addresses the core technical challenge of the JD, despite her profile…

Transferable skill mappings:

• Migration of keyword-search to embedding-based retrieval → Embeddings-based retrieval systems
• Offline-to-online correlation and evaluation rigor → Evaluation frameworks for ranking systems
• Terraform and MLOps production experience → MLOps

Aarohi Sen (Rank 25 · Composite 74.6)

Aarohi is a classic 'infrastructure-first' practitioner who likely ranked lower due to a lack of buzzword-heavy distributed systems terminology in her profile. However, her hands-on experience building production-grade semantic search and ranking systems directly satisfies the core technical requirements of the JD. Her ability to bridge the gap between raw infrastructure (Kubeflow/MLflow) and business outcomes (12% revenue lift) suggests she…

Transferable skill mappings:

• Built semantic search infrastructure using sentence-transformers and FAISS → Embeddings-based retrie
• Built and operated production ML pipelines using MLflow and Kubeflow → MLOps
• Evolved e-commerce search ranking from hand-tuned to learning-to-rank → Evaluation frameworks for ra

Methodology

Composite Score Formula

The composite score (0–100) is computed from four evidence-based dimensions using the following canonical formula defined in SKILLS.md:

```
composite_score = (fit_score × 0.35) + (impact_score × 0.30) + (potential_score × 0.20) + ((100 − risk_score) × 0.15)
```

fit_score (35%)
Measures alignment between the candidate's background and the JD's explicit and implicit requirements. Driven by career history title classification, ML keyword evidence in job descriptions, and product vs. services company history.

impact_score (30%)
Measures real-world measurable outcomes: quantified achievements (numbers, scale, revenue, users). A candidate with zero quantified impact signals scores at most 40 on this dimension — missing numbers cannot be assumed.

potential_score (20%)
Measures growth trajectory: career velocity (promotions per year), complexity growth across chronological career history, and self-learning signals such as OSS contributions, certifications, and hackathons.

risk_score (15%, inverted)
Measures hiring risk: skill gaps vs. JD must-haves, very short tenures, no evidence of collaboration, overqualification signals, or pure-services career history. Higher risk = lower composite via the (100 – risk_score) inversion.

Evidence-Based Scoring

All scores are derived from evidence traceable to the candidate's actual profile text. The system never invents skills, impact numbers, or experience not present in the source data. Each score includes a confidence_level field (high / medium / low) indicating how much direct textual evidence supports it. Low-confidence scores default to 50 and are flagged for human review.

Honeypot Detection

The dataset contains candidates with logically impossible profiles designed to trap keyword-matching systems. This pipeline detects and scores them 0.0 using four rules: impossible experience timelines, mass expert skills with zero usage duration, copy-pasted role descriptions, and high profile completeness with empty career content. 472 honeypot candidates were detected and excluded from the shortlist.